

# CSGI-GA 3303-076 Vision meets ML

## Homework 2

Enter your name and NetID below.

**Name: Shuvadeep Saha**

**NetID: ss15592**

This assignment has **2 PARTS!**

The main goals of this assignment include:

Part1:

- 1. Giving an introduction to PyTorch
- 2. Loading images using PyTorch DataLoader
- 3. Performing Inference on a Pretrained Resnet-18 model using a small dataset
- 4. Doing some analysis on the classification results

Part2:

- 1. Reconstruct a pytorch model from an onnx file
- 2. Build synthetic datasets

Also accompanying parts 3-4, there are a few questions (**8 questions total in part 1, 2 total questions in part 2**). Please give your answers in the space provided. In most of the questions, you will be asked to complete a code snippet. You can quickly navigate to those questions by searching (Ctrl/Cmd-F) for `TODO`: . For part 2, please append your code to the same Jupyter notebook just like HW1. You are only required to submit a PDF of your notebook with the required outputs. Please note that we will not go through ipynbs and also make sure all the outputs are clear and present.

In this homework, we'll be covering few basic PyTorch concepts. More will be covered in the next homework.

### Part 1: Introduction to PyTorch

[PyTorch \(https://pytorch.org/\)](https://pytorch.org/) and [TensorFlow \(https://www.tensorflow.org/\)](https://www.tensorflow.org/) are open source machine learning frameworks that are used to develop and train neural network models. PyTorch is primarily developed by Facebook AI Research lab (FAIR), whereas TensorFlow is developed by the Google Brain team. You can read more about the differences between PyTorch and Tensorflow from these 2 articles:

- 1. [Real Python blog post by Ray Johns, "PyTorch vs TensorFlow for Your Python Deep Learning Project" \(https://realpython.com/pytorch-vs-tensorflow/\)](https://realpython.com/pytorch-vs-tensorflow/)
- 2. [Towards Data Science article by Kirill Dubovikov, "PyTorch vs TensorFlow — spotting the difference" \(https://towardsdatascience.com/pytorch-vs-tensorflow-spotting-the-difference-25c75777377b\)](https://towardsdatascience.com/pytorch-vs-tensorflow-spotting-the-difference-25c75777377b)

In this course, we will use PyTorch.

A neural network consists of several layers. Each layer takes an input array, computes a linear product, applies some non linear activation function to give an output. These inputs and outputs are just n-dimensional arrays, and are called as `Tensor` s in PyTorch. These are similar to what we have in `NumPy` except that `Tensor` s can run on GPUs. PyTorch provides several functions for operating on these `Tensor` s. In the background, during forward and backward propagation, `Tensor` s can keep track of the computational graph and its weights' gradients.

We'll now show how one can build a neural network in PyTorch using a simple example. [Credits: [Deep Learning with PyTorch: A 60 Minute Blitz by Soumith Chintala \(https://pytorch.org/tutorials/beginner/blitz/neural\\_networks\\_tutorial.html#neural-networks\)](https://pytorch.org/tutorials/beginner/blitz/neural_networks_tutorial.html#neural-networks)]

Consider this network for classifying digit images:



This network can be built in PyTorch using the `torch.nn` package. `torch.nn` implements various neural network layers (full-connected, convolution, pooling, etc.), and activation functions (softmax, sigmoid, etc.).

```
In [1]: import torch
import torch.nn as nn
import torch.nn.functional as F

# nn.Module is a container that's used as the base class for
# building any network. It has built-in methods for loading parameters
# from a pretrained network, exporting weights from a network, moving
# the model from CPU to GPU, etc.

class Net(nn.Module):

    def __init__(self):
        super(Net, self).__init__()
        # Docs:
        # https://pytorch.org/docs/stable/generated/torch.nn.Conv2d.html#torch.nn.Conv2d
        # https://pytorch.org/docs/stable/generated/torch.nn.Linear.html#torch.nn.Linear

        # 1 input image channel, 6 output channels, 3x3 square convolution
        # kernel
        self.conv1 = nn.Conv2d(1, 6, 3)
        self.conv2 = nn.Conv2d(6, 16, 3)

        # an affine operation: y = Wx + b
        self.fc1 = nn.Linear(16 * 6 * 6, 120) # 6*6 from image dimension
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        # Docs:
        # https://pytorch.org/docs/stable/nn.functional.html#max-pool2d
        # https://pytorch.org/docs/stable/tensor_view.html#tensor-view-doc
        # https://pytorch.org/docs/stable/nn.functional.html#relu

        # Max pooling over a (2, 2) window
        x = F.max_pool2d(F.relu(self.conv1(x)), (2, 2))
        # If the size is a square you can only specify a single number
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)
        x = x.view(-1, self.num_flat_features(x))
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x

    def num_flat_features(self, x):
        size = x.size()[1:] # all dimensions except the batch dimension
        num_features = 1
        for s in size:
            num_features *= s
        return num_features
```

nn.Module subclass requires us to fill in 3 methods:

- 1. `__init__` : Calls the base constructor, and instantiates all the network layers (ones that have trainable parameters).
- 2. `forward` : Defines the computation flow for forward propagation/inference by specifying a chain of operations. Given an input Tensor `x` , it computes output from one layer, feeds it to the next, and so on, till the final output [i.e. the digit class (0-9) the image belongs to] is computed.

You can read more from [this PyTorch tutorial on Neural Networks \(https://pytorch.org/tutorials/beginner/blitz/neural\\_networks\\_tutorial.html\)](https://pytorch.org/tutorials/beginner/blitz/neural_networks_tutorial.html). First 2 sections (*Neural Networks* and *Define the network*) are enough. We will not cover backprop/autograd in this homework though.

With the network class defined, let's instantiate the network and print the layers inside it.

```
In [2]: net = Net()
print(net)

Net(
  (conv1): Conv2d(1, 6, kernel_size=(3, 3), stride=(1, 1))
  (conv2): Conv2d(6, 16, kernel_size=(3, 3), stride=(1, 1))
  (fc1): Linear(in_features=576, out_features=120, bias=True)
  (fc2): Linear(in_features=120, out_features=84, bias=True)
  (fc3): Linear(in_features=84, out_features=10, bias=True)
)
```

## Part 2: Loading images using Dataloaders

As you've read in the above PyTorch tutorial, we can compute the output of a network on an image by calling `net()` on its tensor. This will return an output Tensor , with one value for each of the 10 classes. The network's predicted label for this image would be the label corresponding to the index with the highest value.

Consider this example below. An image of size [1, 32, 32] is randomly generated and is fed into the network to get its label. The output is a tensor containing 10 float values.

```
In [3]: # Use a random 32x32 input
# 1st dimension = 1 refers to batch size = 1 (only one example/image in this batch of examples)
# 2nd dimension = 1 refers to the number of input channels = 1 (this is a BW image)
input_tensor = torch.randn(1, 1, 32, 32)

# Do inference on this example
output_tensor = net(input_tensor)
print(output_tensor)

tensor([[ -0.0404,  0.0995,  0.0768,  0.0877,  0.1315, -0.0896,  0.0573,  0.1123,
          -0.0239, -0.1045]], grad_fn=<AddmmBackward0>)
```

```
In [4]: print("The predicted class index of the random image is", output_tensor.argmax().numpy())

# argmax -> returns the index of the maximum value in the output_tensor
# numpy -> converts a PyTorch Tensor into a Numpy array

The predicted class index of the random image is 4
```

However, this approach (of calling the network on the input tensor to get its output) doesn't scale well. Typically, during the training phase, one would shuffle all the images in a dataset, load a batch of images from the shuffle, preprocess all the images in that batch (by applying some transformations such as random cropping, rescaling the image, etc.), and then provide the resulting image tensor to the model. If we simply do all these operations by iterating over the dataset using a `for` loop, these operations can turn out to be inefficient, especially when working on large datasets. For every batch, we have to generate the input tensor, and only after this step is complete, the network will be able to compute the output on this tensor.

To avoid this problem, we use a PyTorch utility called `DataLoader`, which is capable of performing all the above operations using `multiprocessing` workers (i.e. multiple threads). By delegating these tasks to separate workers, these workers can fetch the next batch while training on the previous batch is still in progress, thereby reducing the batch generation bottleneck. In other words, the training process doesn't need to wait for the worker threads to fetch the batch for training. Furthermore, this is necessary because loading all the images in a dataset into memory and then processing them may not always be feasible because of memory constraints.

`DataLoader` takes as input a `Dataset` object, which is a PyTorch abstract class for holding and operating on a dataset. This also provides a homogenous way of working with different datasets.

`torchvision.datasets` (<https://pytorch.org/vision/0.8/datasets.html#>) package provides some commonly used CV-related datasets such as `MNIST` (<http://yann.lecun.com/exdb/mnist/>), `CIFAR10` (<https://www.cs.toronto.edu/~kriz/cifar.html>), `Places365` (<http://places2.csail.mit.edu/index.html>), etc.

Using `Dataset` and `DataLoader`, here's how the dataloading phase in a PyTorch program works:

```
In [5]: from torchvision import datasets, transforms
from torch.utils.data import DataLoader, Subset

# Transforms refer to a set of operations one would apply to each
# image before sending it to the network. For example, all images
# in the dataset may not be of the same size, in which case, we could
# use a transform to scale the image to a given size.
# PyTorch provides a variety of transforms:
# https://pytorch.org/docs/stable/torchvision/transforms.html
transform = transforms.Compose([
    transforms.ToTensor(), # first, convert image to PyTorch tensor
    transforms.Normalize((0.1307,), (0.3081,)) # normalize inputs
])
# We're telling MNIST dataset is in the folder "./mnist"
# If it doesn't exist in the folder, we're allowing the package
# to download it (by setting download=True)
mnist = datasets.MNIST('./mnist', download=True, transform=transform)

# We don't want to use all the 60000 training images in MNIST.
# Instead we'll just use a subset of the training set (64 images)
# to make the computation faster
mnist = Subset(mnist, range(64))

# DataLoader requires a Dataset object.
# Each call to dataloader will return a batch of 16 images
# Batches are generated after shuffling the dataset
dataloader = DataLoader(mnist, batch_size=16, shuffle=True)

# Typically one would want to train using a GPU.
# To be able to do that, torch needs to locate/access the GPU.
# cuda = torch.device("cuda") # use "cuda:0", "cuda:1" instead to use the first and second GPU respectively

for batch_idx, (inputs, targets) in enumerate(dataloader):
    # To move the tensors to the GPU device, one would do
    # inputs = inputs.to(cuda)
    # targets = targets.to(cuda)
    # However, we currently don't have access to a
    # GPU, so leave the above 2 lines commented.
    print("Batch Idx", batch_idx, "; Input size", inputs.shape, "; Target size", targets.shape)

    # Forward propagation logic sits here
```

```
Batch Idx 0 ; Input size torch.Size([16, 1, 28, 28]) ; Target size torch.Size([16])
Batch Idx 1 ; Input size torch.Size([16, 1, 28, 28]) ; Target size torch.Size([16])
Batch Idx 2 ; Input size torch.Size([16, 1, 28, 28]) ; Target size torch.Size([16])
Batch Idx 3 ; Input size torch.Size([16, 1, 28, 28]) ; Target size torch.Size([16])
```

Each batch has 16 images as we've specified in the `DataLoader` . Each target has 16 values, where each value is the index of the true label the image belongs to. The 2nd dimension in the `inputs` tensor refers to the number of input channels, which is 1 in this case, because MNIST contains BW images. 3rd and 4th dimension refer to the image's height and width respectively.

Also check out this [thread \(https://twitter.com/ScottCondron/status/1363494433715552259\)](https://twitter.com/ScottCondron/status/1363494433715552259) [yes, it's a tweet :)] and [this article \(https://stanford.edu/~shervine/blog/pytorch-how-to-generate-data-parallel\)](https://stanford.edu/~shervine/blog/pytorch-how-to-generate-data-parallel).

Part 3: Inference on the sample Imagenet dataset

To demonstrate inference on a pretrained network, we've created a dataset, consisting of 100 images by sampling 5 images from 20 classes in the Imagenet dataset. Imagenet dataset comprises approximately 1 million images and 1000 classes, and in the interest of keeping things simple, we've provided a smaller subset.

Question 1

To work with this dataset, you'll need to create a `Dataset` object, since Imagenet dataset cannot be downloaded directly via the `torchvision.datasets` package. In order to do that, you need to define a class subclassing the abstract `Dataset` class. Any custom `Dataset` class should contain the following:

- 1. `__init__()` method: Constructor method. We've provided the list of attributes the class would need.
- 2. `__len__()` method: Returns the number of samples in the dataset.
- 3. `__getitem__()` method: Given an index, it returns the (input, target) pair at that index.

```
In [6]: import torch
from torch.utils.data import Dataset
import glob
import numpy as np
from PIL import Image

class SampleImagenetDataset(Dataset):
    def __init__(self, root_dir, transform=None):
        '''
        Args:
        root_dir: path where all the 100 jpg images are located
        transform: set of transforms to apply on a dataset sample
        '''

        # TODO: Generate a list containing
        # paths to all the 100 jpg images. Use the glob library
        # to do this.
        self.image_list = glob.glob('/home/jovyan/HW2/imagenet_samples/*.jpg')
        #print('This is in imagelist', self.image_list)

        # TODO: There's a file called 'labels_to_ids.txt' inside the root_dir.
        # Read the file contents to generate a dictionary containing
        # label to idx mapping. The resulting dictionary should have a `str`
        # type key mapped to a `int` type value representing the index
        label_dict = {}
        label_file = open('/home/jovyan/HW2/imagenet_samples/labels_to_ids.txt')
        for line in label_file:
            key, value = line.split()
            label_dict[key] = int(value)
        #print (label_dict)

        self.labels_to_ids = label_dict
        self.transform = transform

    def __len__(self):
        return len(self.image_list)

    def __getitem__(self, idx):
        img_path = self.image_list[idx]

        classname = img_path.split("_", 3)[-1].split(".")[0]
        label = self.labels_to_ids[classname]

        # TODO: Open the image using PIL.Image package
        image = Image.open(img_path)

        if self.transform:
            # TODO: Apply the transform to the `image`
            # and overwrite the output back to the `image` variable
            image = self.transform(image)

        return image, label
```

Question 2

Having defined our custom dataset class, we'll instantiate a `SampleImagenetDataset` class object. And then we'll set up a `DataLoader` object.

```
In [7]: from torchvision import transforms
from torch.utils.data import DataLoader, Subset

data_transform = transforms.Compose([
    transforms.Resize((224,224)),
    transforms.ToTensor(),
    transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
])

sample_dataset = SampleImagenetDataset('../shared/HW2/imagenet_samples', data_transform)

# TODO: Create a `DataLoader` object
# to load images from `sample_dataset` object created above.
# Shuffle the dataset and use a batch size of 10.
# Docs: https://pytorch.org/docs/stable/data.html#torch.utils.data.DataLoader
dataset_loader = DataLoader(sample_dataset, batch_size=10)
```

Question 3

Using the `DataLoader` object, we will iterate over the dataset, dividing it into batches and working on each batch per iteration. For each batch, the network takes the input tensor and outputs the prediction. All predictions and targets (true labels) will be combined and stored in 2 lists -- `outputs` and `targets` respectively.

```
In [8]: from torchvision import models
import numpy as np

# TODO:
# Construct a pretrained ResNet-18 model. Use the
# `torchvision.models` package.
resnet18 = models.resnet18(pretrained=True)

resnet18.eval()
outputs = []
targets = []

for idx, (input_tensor, target_tensor) in enumerate(dataset_loader):
    # TODO: Compute the pretrained model's output on input_tensor
    output_tensor = resnet18(input_tensor)
    #     print('Input', input_tensor.shape)
    #     print('Output', output_tensor.shape)
    #     print('Target', target_tensor.shape)

    # Notes:
    # 1. Using argmax() gives the index of the maximum value in the output_tensor
    #     This is the network's predicted label. output_tensor is a 2d tensor of
    #     dimensions [batch_size, num_classes]. argmax(axis=1) returns the indices
    #     of the maximum values across dim=1 (across the row). Resulting tensor
    #     would be a 1D tensor of size [batch_size].
    #     See https://pytorch.org/docs/stable/generated/torch.argmax.html#torch-argmax
    # 2. Calling numpy() converts a PyTorch tensor to a numpy array
    # 3. list(<arr>) converts a NumPy array to a list
    # 4. l1 += l2 appends elements of list l2 to list l1.
    outputs += list(output_tensor.argmax(1).numpy())
    targets += list(target_tensor.numpy())
```

```
In [9]: # print(torch.argmax(output_tensor[9,:],dim=0))
# print(outputs)
# #output_tensor.size(), input_tensor.size(), idx, len(outputs), targets, target_tensor
```

Question 4

How many batches are generated in the `for` loop? What're the sizes of the `input_tensor` , `target_tensor` and `output_tensor` in any given iteration?

Answer

10 batches are generated in each `for` loop.

For any given iteration, the sizes are as follows:

- `input_tensor` - ([10, 3, 224, 224]); 10 batch-size, 3 channels(RGB), (224,224) image dimensions
- `target_tensor` - ([10]); Refers to the 10 true label index
- `output_tensor` - ([10, 1000]); 1000 predicted label values for each 1000 classes in one row with total of 10 rows (batch-size)

Part 4: Classification Analysis

From part 3, we have the network's predicted label indices and the target label (or the true label) indices. Using [scikit-learn](https://machinelearningmastery.com/a-gentle-introduction-to-scikit-learn-a-python-machine-learning-library/) (<https://machinelearningmastery.com/a-gentle-introduction-to-scikit-learn-a-python-machine-learning-library/>), we'll perform some analysis on those results.

But before we do that, we'll convert those `int` -type indices to `str` -type labels using a `idx-to-label` mapping. This mapping has been provided in the `shared/imagenet_samples` folder in a file called `ids_to_labels.txt` .

Question 5

Read the contents of the file and generate a python dictionary that holds the `idx-to-label` mapping. In the map, the key is an `int` -type index and the value is a `str` -type label.



```
In [10]: # TODO:
ids_to_labels = {}
label_file = open('/home/jovyan/HW2/imagenet_samples/ids_to_labels.txt')
for line in label_file:
    key, value = line.split()
    ids_to_labels[int(key)] = value
print (ids_to_labels)

# We'll convert the list of keys in the mapping to a Python `set`
ids = set(ids_to_labels.keys())

{605: 'iPod', 950: 'orange', 951: 'lemon', 952: 'fig', 795: 'ski', 875: 'trombone', 107: 'jellyfish', 587: 'hammer', 748: 'purse', 957: 'pomegranate', 420: 'banjo', 924: 'guacamole', 802: 'snowmobile', 827: 'stove', 340: 'zebra', 949: 'strawberry', 523: 'crutch', 534: 'dishwasher', 530: 'digital_clock', 543: 'dumbbell'}
```

Now let's convert these `int` -values in the `outputs` list to a list of `str` -type labels. However, some network predictions may not be in the list of classes we've picked (remember: Imagenet has 1000 classes and we just picked 20 classes). Such indices will be mapped to label "other" using the method below:

```
In [11]: def transform_outputs(outputs):
        return [ids_to_labels[output] if output in ids else "other" for output in outputs]

outputs = transform_outputs(outputs)
targets = transform_outputs(targets)

In [12]: # print(outputs)
        # print(targets)
```

Question 6

Using `sklearn.metrics` ([https://scikit-learn.org/stable/modules/model\\_evaluation.html#classification-metrics](https://scikit-learn.org/stable/modules/model_evaluation.html#classification-metrics)) generate a classification report showing the per-class precision , recall and f1-score . Also print the accuracy of the predictions, macro and micro averages of precision, recall and f1-score across all classes .

If the method you're using does not print the micro-averaged metrics, please briefly explain why in 3-4 lines. In that case, also report the values of these 3 metrics.

```
In [13]: from sklearn import metrics

# TODO:
report = metrics.classification_report(targets, outputs, zero_division=1)
print(report)
```

	precision	recall	f1-score	support
banjo	1.00	1.00	1.00	5
crutch	1.00	0.80	0.89	5
digital_clock	1.00	0.60	0.75	5
dishwasher	1.00	0.80	0.89	5
dumbbell	1.00	0.40	0.57	5
fig	1.00	1.00	1.00	5
guacamole	1.00	0.60	0.75	5
hammer	1.00	0.40	0.57	5
iPod	0.67	0.40	0.50	5
jellyfish	1.00	0.60	0.75	5
lemon	1.00	0.80	0.89	5
orange	1.00	1.00	1.00	5
other	0.00	1.00	0.00	0
pomegranate	1.00	1.00	1.00	5
purse	1.00	0.40	0.57	5
ski	1.00	0.80	0.89	5
snowmobile	1.00	0.80	0.89	5
stove	1.00	1.00	1.00	5
strawberry	1.00	1.00	1.00	5
trombone	1.00	0.40	0.57	5
zebra	1.00	0.80	0.89	5
accuracy			0.73	100
macro avg	0.94	0.74	0.78	100
weighted avg	0.98	0.73	0.82	100

```
In [14]: micro_precision = metrics.precision_score(targets, outputs, average='micro')
micro_recall = metrics.recall_score(targets, outputs, average='micro')
micro_f1 = metrics.f1_score(targets, outputs, average='micro')

print('The metrics for micro-averaged are below:')
print('Precision: ', micro_precision)
print('recall: ', micro_recall)
print('f1:', micro_f1)
```

The metrics for micro-averaged are below:  
Precision: 0.73  
recall: 0.73  
f1: 0.7299999999999999

Question 7

In 3-4 sentences, explain what micro- and macro- averages are and how they're different from each other.

Answer

The macro average simply calculates the mean of the metrics, giving equal weight to each class. This is to say that infrequent classes (classes that occur rarely) that might be present in the dataset doesn't affect this average rather macro-averaging may be a means of highlighting their performance.

On the other hand micro gives each sample-class pair an equal contribution and is more suitable to assess the model performance in case where there is an uneven distribution of model classes. This highlights

Since we have equally distributed classes over the sample size, we have taken the macro to see the overall performance of the model.

Question 8

A confusion matrix (or an error matrix) allows us to visualize model's performance on a per-class basis. From [Wikipedia \(https://en.wikipedia.org/wiki/Confusion\\_matrix\)](https://en.wikipedia.org/wiki/Confusion_matrix), "each row of the matrix represents the instances in a predicted class, while each column represents the instances in an actual class (or vice versa)".

Using sklearn.metrics , compute the confusion matrix.

```
In [15]: from sklearn.metrics import confusion_matrix


ids_to_labels[0] = "other"
labels_list = list(ids_to_labels.values())
# TODO:
conf_matrix = confusion_matrix(targets, outputs, labels=labels_list)
```

Now let's plot the confusion matrix:



```
In [16]: import matplotlib.pyplot as plt
%matplotlib inline

fig, ax = plt.subplots(figsize=(7, 7))
plot = ax.imshow(conf_matrix, interpolation="nearest")
ax.set_xticks(np.arange(0, 21))
ax.set_xticklabels(labels_list)
plt.setp(ax.get_xticklabels(), rotation=45, ha="right", rotation_mode="anchor")
ax.set_yticks(np.arange(0, 21))
ax.set_yticklabels(labels_list);
ax.set_xlabel("Predicted class")
ax.set_ylabel("Target class")
plt.colorbar(plot, ax=ax);
```



Please complete part2 after this cell

```
In [1]: import torch
from torch.nn import Linear, ReLU, Conv2d, MaxPool2d, Module, Sequential, AdaptiveAvgPool2d, ConvTranspose2d
from torchsummary import summary
import torch.onnx

class my_CNN(Module):
    def __init__(self):
        super(my_CNN, self).__init__()
        self.block1 = Sequential(
            # Layer 1
            Conv2d(in_channels=3, out_channels=64, kernel_size=3, padding=1, bias=False),
            ReLU(inplace=True),
            MaxPool2d(kernel_size=2, stride=2),
            # Layer 2
            Conv2d(in_channels=64, out_channels=128, kernel_size=3, padding=1, bias=False),
            ReLU(inplace=True),
            MaxPool2d(kernel_size=2, stride=2)
        )

        self.block2 = Sequential(
            # Layer 3
            Conv2d(in_channels=128, out_channels=256, kernel_size=3, padding=1, bias=False),
            ReLU(inplace=True),
            MaxPool2d(kernel_size=2, stride=2)
        )

        self.block3 = Sequential(
            # Layer 4
            Conv2d(in_channels=256, out_channels=512, kernel_size=3, padding=1, bias=False),
            ReLU(inplace=True),
            MaxPool2d(kernel_size=2, stride=2)
        )

        self.block4 = Sequential(
            # Layer 5
            Conv2d(in_channels=512, out_channels=1024, kernel_size=3, padding=1, bias=False),
            ReLU(inplace=True),
            MaxPool2d(kernel_size=2, stride=2),
            # Layer 6
            ConvTranspose2d(in_channels=1024, out_channels=512, kernel_size=2, stride=2, bias=True)
        )

        self.block5 = Sequential(
            # Layer 7
            ConvTranspose2d(in_channels=1024, out_channels=256, kernel_size=2, stride=2, bias=True)
        )

        self.block6 = Sequential(
            # Layer 8
            ConvTranspose2d(in_channels=512, out_channels=128, kernel_size=2, stride=2, bias=True)
        )

        self.block7 = Sequential(
            # Layer 9
            ConvTranspose2d(in_channels=256, out_channels=128, kernel_size=2, stride=2, bias=True)
        )

        self.block8 = Sequential(
            # Layer 10
            ConvTranspose2d(in_channels=128, out_channels=3, kernel_size=2, stride=2, bias=True)
        )

    def forward(self, input_image):
        input_image1 = self.block1(input_image)
        input_image2 = self.block2(input_image1)
        input_image3 = self.block3(input_image2)
        input_image4 = self.block4(input_image3)
        input_image5 = torch.cat((input_image3, input_image4), axis=1)
        input_image5 = self.block5(input_image5)
        input_image6 = torch.cat((input_image2, input_image5), axis=1)
        input_image6 = self.block6(input_image6)
        input_image7 = torch.cat((input_image1, input_image6), axis=1)
        input_image8 = self.block7(input_image7)
        input_image9 = self.block8(input_image8)
        input_image = input_image9
        return input_image

my_torch_model = my_CNN()
inputs = torch.randn(1, 3, 256, 256, requires_grad=True)
torch_out = my_torch_model(inputs)
summary(my_torch_model,(3, 224, 224))
torch.onnx.export(my_torch_model,inputs,"My_CNN_model.onnx", export_params=True, opset_version=10)
```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 64, 224, 224]	1,728
ReLU-2	[-1, 64, 224, 224]	0
MaxPool2d-3	[-1, 64, 112, 112]	0
Conv2d-4	[-1, 128, 112, 112]	73,728
ReLU-5	[-1, 128, 112, 112]	0
MaxPool2d-6	[-1, 128, 56, 56]	0
Conv2d-7	[-1, 256, 56, 56]	294,912
ReLU-8	[-1, 256, 56, 56]	0
MaxPool2d-9	[-1, 256, 28, 28]	0
Conv2d-10	[-1, 512, 28, 28]	1,179,648
ReLU-11	[-1, 512, 28, 28]	0
MaxPool2d-12	[-1, 512, 14, 14]	0
Conv2d-13	[-1, 1024, 14, 14]	4,718,592
ReLU-14	[-1, 1024, 14, 14]	0
MaxPool2d-15	[-1, 1024, 7, 7]	0
ConvTranspose2d-16	[-1, 512, 14, 14]	2,097,664
ConvTranspose2d-17	[-1, 256, 28, 28]	1,048,832
ConvTranspose2d-18	[-1, 128, 56, 56]	262,272
ConvTranspose2d-19	[-1, 128, 112, 112]	131,200
ConvTranspose2d-20	[-1, 3, 224, 224]	1,539
Total params: 9,810,115		
Trainable params: 9,810,115		
Non-trainable params: 0		
Input size (MB): 0.57		
Forward/backward pass size (MB): 125.56		
Params size (MB): 37.42		
Estimated Total Size (MB): 163.56		

In [ ]:

In [1]:

```
import cv2
import numpy as np
from IPython.display import Image
```

In [2]:

```
def rhombus(box_coord, image_width, image_height, w_x, h_y):
    image = np.zeros((image_height, image_width), dtype=int)
    center = [box_coord[0]+w_x//2, box_coord[1]+h_y//2]
    x, y = center[1], center[0] # Rows and columns of the center

    cx,cy = [x, y]
    queue = [[cx,cy]]
    neighbors = [[1, 0], [-1, 0], [0, 1], [0, -1]]

    if w_x < h_y:
        min_ax = w_x//2
        maj_ax = h_y//2
    else:
        min_ax = h_y//2
        maj_ax = w_x//2

    def dist(a,b):
        return (abs(a[0]-b[0])*min_ax) + (abs(a[1]-b[1])*maj_ax)

    while(len(queue)>0):
        tx, ty = queue.pop(0)
        if(dist([cx,cy], [tx,ty])) >= (min_ax*maj_ax) or image[ty,tx]==1:
            continue # Do not need to fill this value, because outside the bounds
        image[ty,tx] = 1 # mark the pixel; we'll multiply the mask by 255 for visualization later

        for i,j in neighbors:
            x = tx + i
            y = ty + j
            if(0<=x<img.shape[0] and 0<=y<img.shape[1]): #boundary check to make sure co-ordinates are within
the image
                queue.append((x,y))

    return np.transpose(image)
```

In [3]:

```
def rect(box_coord, image_width, image_height, w_x, h_y):
    image = np.zeros((image_height, image_width), dtype=int)
    center = [box_coord[0]+w_x//2, box_coord[1]+h_y//2]
    x, y = center[1], center[0] # Rows and columns of the center

    cx,cy = [x, y]
    queue = [[cx,cy]]
    neighbors = [[1, 0], [-1, 0], [0, 1], [0, -1]]

    if w_x < h_y:
        min_ax = w_x//2
        maj_ax = h_y//2
    else:
        min_ax = h_y//2
        maj_ax = w_x//2

    def dist(a,b):
        return (abs((a[1]-b[1])*maj_ax + (a[0]-b[0])*min_ax) + abs((a[1]-b[1])*maj_ax - (a[0]-b[0])*min_ax))

    while(len(queue)>0):
        tx, ty = queue.pop(0)
        if(dist([cx,cy], [tx,ty])) >= (2*min_ax*maj_ax) or image[ty,tx]==1:
            continue # Do not need to fill this value, because outside the bounds
        image[ty,tx] = 1 # mark the pixel; we'll multiply the mask by 255 for visualization later

        for i,j in neighbors:
            x = tx + i
            y = ty + j
            if(0<=x<img.shape[0] and 0<=y<img.shape[1]): #boundary check to make sure co-ordinates are within
the image
                queue.append((x,y))

    return np.transpose(image)
```

```
In [4]: def circle(box_coord, image_width, image_height, w_x, h_y):
        image = np.zeros((image_height, image_width), dtype=int)
        center = [box_coord[0]+w_x//2, box_coord[1]+h_y//2]
        x, y = center[1], center[0] # Rows and columns of the center

        cx,cy = [x, y]
        queue = [[cx,cy]]
        neighbors = [[1, 0], [-1, 0], [0, 1], [0, -1]]
        if w_x < h_y:
            rad_cir = w_x//2
        else:
            rad_cir = h_y//2

        def dist(a,b):
            return np.sqrt((a[0]-b[0])**2 + (a[1]-b[1])**2)

        while(len(queue)>0):
            tx, ty = queue.pop(0)
            if(dist([cx,cy], [tx,ty])) > rad_cir or image[ty,tx]==1:
                continue # Do not need to fill this value, because outside the bounds
            image[ty,tx] = 1 # mark the pixel; we'll multiply the mask by 255 for visualization later

            for i,j in neighbors:
                x = tx + i
                y = ty + j
                if(0<=x<img.shape[0] and 0<=y<img.shape[1]): #boundary check to make sure co-ordinates are within
the image
                    queue.append((x,y))

        return np.transpose(image)
```

```
In [5]: def ellipse(box_coord, image_width, image_height, w_x, h_y):
        image = np.zeros((image_height, image_width), dtype=int)
        center = [box_coord[0]+w_x//2, box_coord[1]+h_y//2]
        x, y = center[1], center[0] # Rows and columns of the center

        cx,cy = [x, y]
        queue = [[cx,cy]]
        neighbors = [[1, 0], [-1, 0], [0, 1], [0, -1]]

        if w_x < h_y:
            min_ax = w_x//2
            maj_ax = h_y//2
        else:
            min_ax = h_y//2
            maj_ax = w_x//2

        def dist(a,b):
            return np.sqrt((((a[0]-b[0])**2)*(min_ax**2)) + (((a[1]-b[1])**2)*(maj_ax**2)))

        while(len(queue)>0):
            tx, ty = queue.pop(0)
            if(dist([cx,cy], [tx,ty])) >= (min_ax*maj_ax) or image[ty,tx]==1:
                continue # Do not need to fill this value, because outside the bounds
            image[ty,tx] = 1 # mark the pixel; we'll multiply the mask by 255 for visualization later

            for i,j in neighbors:
                x = tx + i
                y = ty + j
                if(0<=x<img.shape[0] and 0<=y<img.shape[1]): #boundary check to make sure co-ordinates are within
the image
                    queue.append((x,y))

        return np.transpose(image)
```

```
In [6]: def image_1():
        list_ = open("annotations.txt")
        var = []
        lines = list_.readlines()[8:]
        for i in range(len(lines)):
            temp = lines[i].split()[0:4]
            var.append(list(map(int, temp)))
        return var

def image_2():
    list_ = open("annotations.txt")
    lines = list_.readlines()[2]
    var = list(map(int, lines.split()[0:4]))
    return var

def image_3():
    list_ = open("annotations.txt")
    lines = list_.readlines()[5]
    var = list(map(int, lines.split()[0:4]))
    return var
```

```
In [7]: def loops(var, img, image_width, image_height,i):
        box_coord = [var[0], var[1]]      # The coordinates of the bounding box [x-dir, y-dir]
        w_x = var[2]    # Width of the bouning box
        h_y = var[3]    # Height of the bounding box

        image_rhom = rhombus(box_coord, image_width, image_height, w_x, h_y)
        image_rect = rect(box_coord, image_width, image_height, w_x, h_y)
        image_cir = circle(box_coord, image_width, image_height, w_x, h_y)
        image_ellipse = ellipse(box_coord, image_width, image_height, w_x, h_y)

        return image_rhom, image_rect, image_cir, image_ellipse
```

```
In [8]: def print_to_file(image, c, image_width, image_height, img, string):
        image_ = np.zeros((image_width,image_height,3),dtype= np.uint8)
        image_[:::,0] = image
        image_[:::,1] = image
        image_[:::,2] = image

        img_save = "overlaid_image_"+str(c)+"_"+string+".jpg"
        dst = cv2.addWeighted(img, 0.5, image_*255, 0.5, 0.0)
        cv2.imwrite("/home/jovyan/HW2/"+img_save, dst)
        img_file_path = "/home/jovyan/HW2/"+img_save
        display(Image(img_file_path))
```



```
In [9]: for i in range(3):
        path = '/home/jovyan/HW2/'+str(i+1)+'.jpg'
        img = cv2.imread(path)
        image_width = img.shape[0]
        image_height = img.shape[1]
        image_rhom = np.zeros((image_width,image_height), dtype = int)
        image_rect = np.zeros((image_width,image_height), dtype = int)
        image_cir = np.zeros((image_width,image_height), dtype = int)
        image_ellipse = np.zeros((image_width,image_height), dtype = int)

        if i==0:
            var = image_1()
            for j in range(len(var)):
                image_rhom += loops(var[j], img, image_width, image_height, j)[0]
                image_rect += loops(var[j], img, image_width, image_height, j)[1]
                image_cir += loops(var[j], img, image_width, image_height, j)[2]
                image_ellipse += loops(var[j], img, image_width, image_height, j)[3]
            print_to_file(image_rhom, i+1, image_width, image_height, img, 'rhombus')
            print_to_file(image_rect, i+1, image_width, image_height, img, 'rectangle')
            print_to_file(image_cir, i+1, image_width, image_height, img, 'circle')
            print_to_file(image_ellipse, i+1, image_width, image_height, img, 'ellipse')

        elif i==1:
            var = image_2()
            image_rhom = loops(var, img, image_width, image_height,i+1)[0]
            image_rect = loops(var, img, image_width, image_height,i+1)[1]
            image_cir = loops(var, img, image_width, image_height,i+1)[2]
            image_ellipse = loops(var, img, image_width, image_height,i+1)[3]
            print_to_file(image_rhom, i+1, image_width, image_height, img, 'rhombus')
            print_to_file(image_rect, i+1, image_width, image_height, img, 'rectangle')
            print_to_file(image_cir, i+1, image_width, image_height, img, 'circle')
            print_to_file(image_ellipse, i+1, image_width, image_height, img, 'ellipse')

        elif i==2:
            var = image_3()
            image_rhom = loops(var, img, image_width, image_height,i+1)[0]
            image_rect = loops(var, img, image_width, image_height,i+1)[1]
            image_cir = loops(var, img, image_width, image_height,i+1)[2]
            image_ellipse = loops(var, img, image_width, image_height,i+1)[3]
            print_to_file(image_rhom, i+1, image_width, image_height, img, 'rhombus')
            print_to_file(image_rect, i+1, image_width, image_height, img, 'rectangle')
            print_to_file(image_cir, i+1, image_width, image_height, img, 'circle')
            print_to_file(image_ellipse, i+1, image_width, image_height, img, 'ellipse')
```









*Central Lyon MS Band  
Emily Bisbee, Direrctor  
Rock Rapids, Iowa  
Sketch #3B  
Sample Sketch*

Official Sponsor

This sketch designed by and property of

**DeMOULIN**  
BROS. & CO., Greenville, Illinois 62246

Chris McCurdy  
Area Representative  
877-389-1429  
cmccurdy@demoulin.com

We want to be YOUR uniform company

*Central Lyon MS Band  
Emily Bisbee, Direrctor  
Rock Rapids, Iowa  
Sketch #3B  
Sample Sketch*

Official Sponsor

This sketch designed by and property of

**DeMOULIN**  
BROS. & CO., Greenville, Illinois 62246

Chris McCurdy  
Area Representative  
877-389-1429  
cmccurdy@demoulin.com

We want to be YOUR uniform company



*Central Lyon MS Band  
Emily Bisbee, Direrctor  
Rock Rapids, Iowa  
Sketch #3B  
Sample Sketch*

Official Sponsor

This sketch designed by and property of

**DeMOULIN**  
BROS. & CO., Greenville, Illinois 62246

We want to be YOUR uniform company

Chris McCurdy  
Area Representative  
877-389-1429  
cmccurdy@demoulin.com



*Central Lyon MS Band  
Emily Bisbee, Direrctor  
Rock Rapids, Iowa  
Sketch #3B  
Sample Sketch*





Official Sponsor



Chris McCurdy  
Area Representative  
877-389-1429  
[cmccurdy@demoulin.com](mailto:cmccurdy@demoulin.com)

This sketch designed by and property of



**DeMOULIN**  
BROS. & CO., Greenville, Illinois 62246

We want to be YOUR uniform company















In [ ]: