

# CSGI-GA-3303076 Vision meets ML

## Homework 3 Part 1

Enter your name and NetID below.

Name: Shuvadeep Saha

NetID: ss15592

The main goals of this assignment include:

- 1. Giving an introduction to Mask-RCNN
- 2. Training the predictors for a given dataset
- 3. Finetuning the entire network for the same dataset
- 4. Building an entire segmentation pipeline

There are **TWO PARTS** in this assignment.

Part1: Accompanying each part, there are a few questions (**12 questions in total**) -- 10 mandatory + 2 extra credit. The first 10 questions are worth 100 points and the extra credit questions are worth 20 points.

Part2: Please use HW3Part2.ipynb for part 2. Merge the pdf exports of both notebook and submit **1 PDF**.


Please give your answers in the space provided. This homework has a mix of conceptual and coding questions. You can quickly navigate to coding questions by searching (Ctrl/Cmd-F) for `TODO`: .

### Part 1: Introduction to Mask-RCNN

[Mask-RCNN \(https://arxiv.org/pdf/1703.06870.pdf\)](https://arxiv.org/pdf/1703.06870.pdf) is a network used for instance segmentation. Instance segmentation can be thought of as a hybrid of semantic segmentation and object detection. In other words, we don't want to just find the bounding boxes for each object in our image, we're also interested in finding the segmentation mask of *each object instance*.



Object Detection



Semantic Segmentation



Instance Segmentation

Image Credits: <https://towardsdatascience.com/single-stage-instance-segmentation-a-review-1eeb66e0cc49>  
(<https://towardsdatascience.com/single-stage-instance-segmentation-a-review-1eeb66e0cc49>)

Mask-RCNN is built on top of Faster-RCNN, which is a network used for object detection. Faster-RCNN has 2 outputs for each candidate object (Region of Interest or RoI) - a class label and a bounding box offset. Mask-RCNN adds a third branch to Faster-RCNN for predicting segmentation masks on each RoI.

We'll first briefly go over Faster-RCNN. Faster-RCNN has 2 stages:

- 1. **Region Proposal Network (RPN):** Given the image, it proposes candidate object bounding boxes. Previous object detection models such as RCNN and Fast-RCNN handled this separately from the CNN model. Faster-RCNN takes a different approach -- it integrates these two components into the same network to achieve speedup.
- 2. **Fast-RCNN:** This stages takes each candidate RoI and extracts features from the image feature vector using RoIPool. Using these RoI features, it performs classification and bounding box regression.




Figure 1. The **Mask R-CNN** framework for instance segmentation.

Mask-RCNN has the same 2-stage procedure, but in the 2nd stage, instead of just predicting the classification label and the bounding box offset, it predicts **in parallel** a binary mask for each RoI.

Mask-RCNN relies on a pretrained network (called the "backbone" in the paper) to extract features from the image. These features are fed into the Region Proposal Network (RPN) to generate candidate Rols. For each RoI candidate, a fixed size RoI feature vector is generated using an RoIAlign layer. This RoI feature map is then provided to the classifier, bounding box predictor and the segmentation mask to generate the final output.

Question 1

Training uses a multi-task loss function. What are the three components in this loss function? Is the loss computed per image or per RoI?

Answer:

The multi-class loss function is defined, during training on each RoI. The three components of the loss function are: Classification loss, Bounding-box regression loss and average binary cross entropy loss.

Question 2

[This blog post \(https://lilianweng.github.io/lil-log/2017/12/31/object-recognition-for-dummies-part-3.html\)](https://lilianweng.github.io/lil-log/2017/12/31/object-recognition-for-dummies-part-3.html) by Lilian Weng gives a nice overview of the object detection (RCNN type) networks. In the blog post, it is mentioned that Mask-RCNN uses RoIAlign instead of RoIPool. Explain briefly in 3-4 lines why this is being done.

Answer

RoIPool first quantizes a floating-number RoI to the discrete granularity of the feature map, this quantized RoI is then subdivided into spatial bins which are themselves quantized, and finally feature values covered by each bin are aggregated (usually by max pooling). Though this does not affect the classification, which is robust to small translations, it has a large negative effect on predicting pixel-accurate masks.

To avoid this, Mask-RCNN uses RoIAlign which use bi-linear interpolation to compute the exact values of the input features at four regularly sampled locations in each RoI bin, and aggregate the result using max or average pooling. This gives pixel-accurate masks for image segmentation.

Question 3

What are the different backbones explored in the Mask-RCNN paper? They are denoted in the paper using network-depth-features nomenclature. What is the advantage of using a ResNet-FPN backbone over a ResNet-C4 backbone for feature extraction?

Answer

The authors evaluated ResNet, ResNeXt networks of depth 50 or 101 layers and FPN (Feature Pyramid Network) as the backbones in the paper.

Using a ResNet-FPN backbone for feature extraction with Mask RCNN gives excellent gains in both accuracy and speed.

Part 2: Training the Predictors for a New Dataset

In this section, we'll start with a pretrained Mask-RCNN model that uses Resnet-50-FPN as the backbone. This model was trained on MS-COCO dataset which is widely used for multiple vision tasks such as object detection, instance segmentation, etc.

MS-COCO has 91 classes (90 for objects + 1 for background). Some sample objects in the dataset include `person` , `car` , `bicycle` , `knife` , `train` , etc.

Along with this homework file, we have also provided another sample dataset (we'll refer to it as the [Nature dataset](https://towardsdatascience.com/custom-instance-segmentation-training-with-7-lines-of-code-ff340851e99b) (<https://towardsdatascience.com/custom-instance-segmentation-training-with-7-lines-of-code-ff340851e99b>)). It isn't a standard dataset, but it's small enough (600 train + 200 test images) and allows us to easily demo finetuning a pretrained Mask-RCNN model. This dataset contains only 2 classes - `squirrel` and `butterfly` .

Our goal in part2 and part3 of this assignment is to take the pretrained Mask-RCNN model and finetune/train it for this dataset. However, here in part2, instead of finetuning the entire network, we'll train only the final layers.

In Homework 2, we've shown how one could feed data into the network using `Dataset` s and `DataLoader` s. We'll use the same strategy here for finetuning the model.

We've based this homework on this [PyTorch tutorial on Object Detection Finetuning](https://pytorch.org/tutorials/intermediate/torchvision_tutorial.html) ([https://pytorch.org/tutorials/intermediate/torchvision\\_tutorial.html](https://pytorch.org/tutorials/intermediate/torchvision_tutorial.html)).

In [1]:

```
1 import numpy as np
2 import torch
3 import torchvision
4
5 import json # for reading from json file
6 import glob # for listing files inside a folder
7 from PIL import Image, ImageDraw # for reading images and drawing masks on them.
8
9
10 # Create a custom dataset class for Nature
11 # dataset subclassing PyTorch's Dataset class
12 class NatureDataset(torch.utils.data.Dataset):
13     def __init__(self, root, transforms):
14         self.transforms = transforms
15
16         # Load all image files, sorting them to
17         # ensure that they are aligned with json files
18         imgs = glob.glob(root + '/*.jpg')
19         imgs += glob.glob(root + '/*.png') # some images are in png format
20         self.imgs = sorted(imgs)
21
22         # Mask data is stored in a json file
23         masks = glob.glob(root + '/*.json')
24         self.masks = sorted(masks)
25
26         # Each image can have multiple object instances, and each
27         # instance is associated with either of these 2 labels.
28
29         # Need to convert str-labels to ids. So we'll use
30         # this label-to-index mapping.
31         # Note: we can't start from 0 because 0 is restricted
32         # to the "background" class
33         self.label_to_id = {'squirrel': 1, 'butterfly': 2}
34
35
36     def __getitem__(self, idx):
37         # Have already aligned images and JSON files; can now
38         # simply use the index to access both images and masks
39         img_path = self.imgs[idx]
40         mask_path = self.masks[idx]
41
42         # Read image using PIL.Image and convert it to an RGB image
43         img = Image.open(img_path).convert("RGB")
44
45         # TODO: Read image height, width and mask data from
46         # the JSON file
47         with open(mask_path, 'r') as fp:
48             # TODO: Using json library read the dictionary
49             # from the fp
50             json_dict = json.load(fp)
51
52             # TODO:
53             height = json_dict['imageHeight']
54
55             # TODO:
56             width = json_dict['imageWidth']
57
58             # TODO:
59             poly_shapes_data = json_dict["shapes"]
60
61         # TODO: Each image can have multiple mask instances.
62         # Using the polygon points, generate the 2d-mask
63         # using PIL's ImageDraw.polygon
64         masks = []
65         labels = []
66         for shape_data in poly_shapes_data:
67             polygon_points = [tuple(point) for point in shape_data['points']]
68
69             # TODO: Using Image.new() create an image of size (width, height)
70             # and fill it with 0s.
71             mask_img = Image.new(mode="L", size=(width, height))
72
73             # TODO: Draw the mask on the base image we just created
74             ImageDraw.Draw(mask_img).polygon(polygon_points, outline=1, fill=1)
75
76             mask = np.array(mask_img)
77             masks.append(mask)
78
79             label = shape_data['label']
80             labels.append(label)
81
82
83         # Each mask instance also has an associated label which is str-type
84         # Convert the str into an int using the mapping we created in __init__
85         labels = [self.label_to_id[label] for label in labels]
86
87         # TODO: Generate the bounding boxes for each instance
88         # from the 2d masks
89         num_objs = len(masks)
90
```

```
91     boxes = []
92     for i in range(num_objs):
93         # TODO: Use np.where() to find where mask[i] == True.
94         # pos will be a 2d-list of indices
95         pos = np.where(masks[i]==True)
96
97         # In pos, find the min x- and y- indices;
98         # max x- and y- indices. This will give us our box bounds.
99
100        # TODO:
101        xmin = min(pos[1])
102
103        # TODO:
104        xmax = max(pos[1])
105
106        # TODO:
107        ymin = min(pos[0])
108
109        # TODO:
110        ymax = max(pos[0])
111
112        boxes.append([xmin, ymin, xmax, ymax])
113
114
115        # Convert everything into a torch.Tensor
116        boxes = torch.as_tensor(boxes, dtype=torch.float32)
117        labels = torch.as_tensor(labels, dtype=torch.int64)
118        masks = torch.as_tensor(masks, dtype=torch.uint8)
119
120        image_id = torch.tensor([idx])
121        area = (boxes[:, 3] - boxes[:, 1]) * (boxes[:, 2] - boxes[:, 0])
122
123        # Assume all instances are not crowd
124        iscrowd = torch.zeros((num_objs,), dtype=torch.int64)
125
126        target = {}
127        target["boxes"] = boxes
128        target["labels"] = labels
129        target["masks"] = masks
130        target["image_id"] = image_id
131        target["area"] = area
132        target["iscrowd"] = iscrowd
133
134        # Apply transforms
135        if self.transforms is not None:
136            img, target = self.transforms(img, target)
137
138        return img, target
139
140
141    def __len__(self):
142        return len(self.imgs)
```

/opt/conda/envs/py37/lib/python3.7/site-packages/tqdm/auto.py:22: TqdmWarning: IPProgress not found. Please update jupyter and ipywidgets. See [https://ipywidgets.readthedocs.io/en/stable/user\\_install.html](https://ipywidgets.readthedocs.io/en/stable/user_install.html) ([https://ipywidgets.readthedocs.io/en/stable/user\\_install.html](https://ipywidgets.readthedocs.io/en/stable/user_install.html))

```
from .autonotebook import tqdm as notebook_tqdm
```

Question 4

Complete the TODO sections in the above block. Specifically:

- 1. Read image height, width and polygon shapes data from the JSON file.
- 2. Generate 2D masks from the polygon points. You can follow this idea: <https://stackoverflow.com/a/3732128> (<https://stackoverflow.com/a/3732128>)
- 3. Generate the bounding boxes from the mask data. Assume that the bounding box is a rectangle with the smallest area enclosing the mask.

You may find this json schema useful:



```
{
  "shapes": [ # list of object instances; masks are represented as polygons
    ## data for instance1
    {
      "label": "" # label for instance1
      "points": [] # 2d list of polygon points [(x1, y1), (x2, y2), ..]
    },
    ## data for instance2
    {
      "label": []
      "points": []
    },
    ..
    ..
  ],
  "imagePath": ""
  "imageData" : ""
  "imageHeight": <integer>
  "imageWidth": <integer>
}
```

In [2]:

```
1 # Alternatively, you could also uncomment the line below to see a sample json file
2 #!cat '../shared/datasets/nature-dataset/train/s (3).json'
```

Having implemented our `NatureDataset` class, let's create the `Dataset` and the `DataLoader` objects. Note that we're not using `torchvision.transforms` , instead we're using transforms provided in a separate script in this directory called `transforms.py` .

In [3]:

```
1 import transforms as T
2
3 def get_transform(train):
4     transforms = []
5     transforms.append(T.ToTensor())
6     if train:
7         transforms.append(T.RandomHorizontalFlip(0.5))
8     return T.Compose(transforms)
9
10 # use our dataset and defined transformations
11 dataset = NatureDataset('../shared/datasets/nature/train', get_transform(train=True))
12 dataset_test = NatureDataset('../shared/datasets/nature/test', get_transform(train=False))
13
14 import utils
15
16 # define training and validation data loaders
17 data_loader = torch.utils.data.DataLoader(
18     dataset, batch_size=2, shuffle=True,
19     collate_fn=utils.collate_fn)
20
21 data_loader_test = torch.utils.data.DataLoader(
22     dataset_test, batch_size=1, shuffle=False,
23     collate_fn=utils.collate_fn)
```

Now let's visualize one image from our dataset.

In [4]:


```
1 import matplotlib.pyplot as plt
2
3 img, targets = dataset[506]
4
5 # np.transpose docs: https://numpy.org/doc/stable/reference/generated/numpy.transpose.html
6 # img is a PyTorch tensor, can convert it to a NumPy tensor by calling .numpy() on it.
7 plt.imshow(np.transpose(img.numpy(), (1, 2, 0)));
```

/opt/conda/envs/py37/lib/python3.7/site-packages/ipykernel\_launcher.py:118: UserWarning: Creating a tensor from a list of numpy.ndarrays is extremely slow. Please consider converting the list to a single numpy.ndarray with `numpy.array()` before converting to a tensor. (Triggered internally at `../torch/csrc/autograd/tensor_new.cpp:201.`)



And here's the corresponding mask.

```
In [5]: 1 plt.imshow(np.transpose(targets['masks'].numpy(), (1, 2, 0)), interpolation='none');
```



We'll be training the final layers of the pretrained Mask-RCNN model (with Resnet-50-FPN backbone) available in the `torchvision` package. So let's download the model:

```
In [6]: 1 model_p2 = torchvision.models.detection.maskrcnn_resnet50_fpn(pretrained=True)
```

Question 5

Recall what we said earlier: For training for the new dataset, we need to modify its box predictor (`FastRCNNPredictor`) and its mask predictor (`MaskRCNNPredictor`) to match with the new dataset. Complete the code cell below.

You may find these docs for `FastRCNNPredictor` and `MaskRCNNPredictor` useful:

```
In [154]: FastRCNNPredictor?
Init signature: FastRCNNPredictor(in_channels, num_classes)
Docstring:
Standard classification + bounding box regression layers
for Fast R-CNN.

Args:
    in_channels (int): number of input channels
    num_classes (int): number of output classes (including background)
Init docstring: Initializes internal Module state, shared by both nn.Module and ScriptModule.
File: /usr/local/lib/python3.9/site-packages/torchvision/models/detection/faster_rcnn.py
Type: type
Subclasses:
```

```
In [155]: MaskRCNNPredictor?
Init signature: MaskRCNNPredictor(in_channels, dim_reduced, num_classes)
Docstring:
A sequential container.
Modules will be added to it in the order they are passed in the constructor.
Alternatively, an ordered dict of modules can also be passed in.

To make it easier to understand, here is a small example::

    # Example of using Sequential
    model = nn.Sequential(
        nn.Conv2d(1,20,5),
        nn.ReLU(),
        nn.Conv2d(20,64,5),
        nn.ReLU()
    )

    # Example of using Sequential with OrderedDict
    model = nn.Sequential(OrderedDict([
        ('conv1', nn.Conv2d(1,20,5)),
        ('relu1', nn.ReLU()),
        ('conv2', nn.Conv2d(20,64,5)),
        ('relu2', nn.ReLU())
    ]))
Init docstring: Initializes internal Module state, shared by both nn.Module and ScriptModule.
File: /usr/local/lib/python3.9/site-packages/torchvision/models/detection/mask_rcnn.py
Type: type
Subclasses:
```

```
In [7]: 1 from torchvision.models.detection.faster_rcnn import FastRCNNPredictor
2 from torchvision.models.detection.mask_rcnn import MaskRCNNPredictor
3
4 # Our new dataset has 3 classes: butterfly, squirrel and background
5 num_classes = 3
6
7 # Get number of input features for the classifier
8 in_features = model_p2.roi_heads.box_predictor.cls_score.in_features
9
10 # TODO: replace the pre-trained head with a new one
11 model_p2.roi_heads.box_predictor = FastRCNNPredictor(in_features, num_classes)
12
13 # Get number of input features for the mask predictor
14 in_features_mask = model_p2.roi_heads.mask_predictor.conv5_mask.in_channels
15 hidden_layer = 256
16
17 # TODO: replace the mask predictor with a new one
18 model_p2.roi_heads.mask_predictor = MaskRCNNPredictor(in_features_mask, hidden_layer, num_classes)
```

Training these layers can take several minutes on a CPU, so we've provided GPUs to make this faster. It should take ~5 mins to run the training in Part2. But before that, we want to ensure that PyTorch is able to access the GPU by printing the device PyTorch is currently (prints `cuda` if it's using a GPU, otherwise it prints `cpu` ).

Now we want to freeze all the layers below these predictors. We can do this by setting the `.requires_grad` attribute of the parameters we want to freeze to `False` . Read more about `requires_grad` from [this PyTorch page on Autograd mechanics](https://pytorch.org/docs/stable/notes/autograd.html#excluding-subgraphs-from-backward) (<https://pytorch.org/docs/stable/notes/autograd.html#excluding-subgraphs-from-backward>).

In order to do the computation on a GPU, we have to move the model from main memory to GPU memory. This can be done by simply calling `.to(device)` on the model. See [the docs](https://pytorch.org/docs/stable/generated/torch.nn.Module.html#torch.nn.Module.to) (<https://pytorch.org/docs/stable/generated/torch.nn.Module.html#torch.nn.Module.to>) for more information.

```
In [8]: 1 import itertools
2
3 # Freeze model and move it to device
4 device = torch.device('cuda') if torch.cuda.is_available() else torch.device('cpu')
5 print("Using", device)
6
7 for param in model_p2.parameters():
8     param.requires_grad = False
9
10 pred_params = itertools.chain(
11     model_p2.roi_heads.mask_predictor.parameters(),
12     model_p2.roi_heads.box_predictor.parameters()
13 )
14
15 for param in pred_params:
16     param.requires_grad = True
17
18 model_p2 = model_p2.to(device)
```

Using cuda

We were able to verify that PyTorch is able to access a GPU. Now let's see the layers inside the `model_p2.roi_heads` to understand what we have modified here (we just modified `box_predictor` and `mask_predictor` ). You could also verify the output below from figure 4 in the Mask-RCNN paper. You'll notice that it's the exact same network on the right part of that figure.



```
In [9]: 1 model_p2.roi_heads
Out[9]: RoIHeads(
  (box_roi_pool): MultiScaleRoIAlign(featmap_names=['0', '1', '2', '3'], output_size=(7, 7), sampling_ratio=2)
  (box_head): TwoMLPHead(
    (fc6): Linear(in_features=12544, out_features=1024, bias=True)
    (fc7): Linear(in_features=1024, out_features=1024, bias=True)
  )
  (box_predictor): FastRCNNPredictor(
    (cls_score): Linear(in_features=1024, out_features=3, bias=True)
    (bbox_pred): Linear(in_features=1024, out_features=12, bias=True)
  )
  (mask_roi_pool): MultiScaleRoIAlign(featmap_names=['0', '1', '2', '3'], output_size=(14, 14), sampling_ratio=2)
  (mask_head): MaskRCNNHeads(
    (mask_fcn1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (relu1): ReLU(inplace=True)
    (mask_fcn2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (relu2): ReLU(inplace=True)
    (mask_fcn3): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (relu3): ReLU(inplace=True)
    (mask_fcn4): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (relu4): ReLU(inplace=True)
  )
  (mask_predictor): MaskRCNNPredictor(
    (conv5_mask): ConvTranspose2d(256, 256, kernel_size=(2, 2), stride=(2, 2))
    (relu): ReLU(inplace=True)
    (mask_fcn_logits): Conv2d(256, 3, kernel_size=(1, 1), stride=(1, 1))
  )
)
```

Question 6

We just printed the head architecture. From the above output, list all the layers we're training along their names.

For example, if we're training mask\_fcn1 of mask\_head you'll specify:

mask\_head.mask\_fcn1 : 2d-Conv layer with 256 input channels, 256 output channels and kernel size = (2, 2).

Note: We're **only** asking for layers with trainable parameters.

Answer

box\_predictor.cls\_score : Linear layer with 1024 input features, 3 output features and bias=True

box\_predictor.bbox\_pred : Linear layer with 1024 input features, 12 output features and bias=True

mask\_predictor.mask\_fcn\_logits : Conv2d layer with 256 input channels, 3 output channels, kernel\_size=(1, 1), stride=(1, 1))

Both the dataloaders and model have been prepared for training. All that remains is to set an optimizer and a learning rate scheduler. When we create an optimizer, we have to provide it the list of trainable parameters.

```
In [10]: 1 # Declare optimizer and lr
2 params = [p for p in model_p2.parameters() if p.requires_grad]
3 optimizer = torch.optim.SGD(params, lr=0.005, momentum=0.9, weight_decay=0.0005)
4 lr_scheduler = torch.optim.lr_scheduler.StepLR(optimizer,
5                                               step_size=3,
6                                               gamma=0.1)
```

Question 7

How many trainable parameters are we passing to the optimizer?

Here's an example to calculate # of trainable parameters in a fully connected layer with:

- 1. an additive bias
- 2. in\_channels = 1024
- 3. out\_channels = 10

The number of trainable parameters here will be 1024\*10 + 10 = 10250.

Answer

The number of trainable parameters are

(256 \* 3) + (1024 \* 3) + (1024 \* 12) + (256 \* 256) = 81664

Now we can finally start the training process.


```
In [11]: 1 num_epochs = 3
2
3 from engine import train_one_epoch, evaluate
4
5 for epoch in range(num_epochs):
6     print("Epoch", epoch)
7     # train for one epoch, printing every 10 iterations
8     train_one_epoch(model_p2, optimizer, data_loader, device, epoch, print_freq=10)
9     # update the learning rate
10    lr_scheduler.step()
11    # evaluate on the test dataset
12    evaluate(model_p2, data_loader_test, device=device)
```

Epoch: [1] [290/300] eta: 0:00:02 lr: 0.005000 loss: 0.3252 (0.3891) loss\_classifier: 0.0354 (0.0410) loss\_box\_reg: 0.0520 (0.0712) loss\_mask: 0.2051 (0.2570) loss\_objectness: 0.0056 (0.0057) loss\_rpn\_box\_reg: 0.0111 (0.0136) time: 0.2901 data: 0.0214 max mem: 1593  
Epoch: [2] [299/300] eta: 0:00:00 lr: 0.005000 loss: 0.3260 (0.3880) loss\_classifier: 0.0359 (0.0417) loss\_box\_reg: 0.0547 (0.0713) loss\_mask: 0.2018 (0.2559) loss\_objectness: 0.0036 (0.0057) loss\_rpn\_box\_reg: 0.0088 (0.0134) time: 0.2979 data: 0.0200 max mem: 1593  
Epoch: [2] Total time: 0:01:28 (0.2952 s / it)  
creating index...  
index created!  
Test: [ 0/200] eta: 0:00:35 model\_time: 0.1653 (0.1653) evaluator\_time: 0.0038 (0.0038) time: 0.1765 data: 0.0071 max mem: 1593  
Test: [100/200] eta: 0:00:13 model\_time: 0.1321 (0.1238) evaluator\_time: 0.0042 (0.0041) time: 0.1440 data: 0.0110 max mem: 1593  
Test: [199/200] eta: 0:00:00 model\_time: 0.1254 (0.1247) evaluator\_time: 0.0037 (0.0039) time: 0.1390 data: 0.0103 max mem: 1593  
Test: Total time: 0:00:27 (0.1392 s / it)  
Averaged stats: model\_time: 0.1254 (0.1247) evaluator\_time: 0.0037 (0.0039)  
Accumulating evaluation results...  
DONE (t=0.04s).  
Accumulating evaluation results...  
DONE (t=0.04s).

In order to analyze the output, you'll need to know what an IoU score is. You can read this blog:  
<https://www.pyimagesearch.com/2016/11/07/intersection-over-union-iou-for-object-detection/>  
(<https://www.pyimagesearch.com/2016/11/07/intersection-over-union-iou-for-object-detection/>)

Question 8

In an image (grey box) of size 12x16, there is an object whose ground truth mask is the box with the dashed edge (green color) and the model's predicted mask is the box with the dash-dotted edge (red color). Calculate the IoU score for this prediction.



Answer

$$IoU = \frac{Intersection}{Union} = \frac{5 * 4}{(5 * 7) + (1 * 5)} = \frac{20}{40} = 0.5$$

Let's see the model's prediction for a sample from the test set.

```
In [12]: 1 img, target = dataset_test[1]
2
3 # put the model in evaluation mode
4 model_p2.eval()
5
6 with torch.no_grad():
7     prediction = model_p2([img.to(device)])
```

Here's the sample image.

```
In [13]: 1 Image.fromarray(img.mul(255).permute(1, 2, 0).byte().numpy())
```



And here's the mask generated by the finetuned model.

```
In [14]: 1 Image.fromarray(prediction[0][ 'mask' ][0, 0].mul(255).byte().cpu().numpy())
```



### Part 3: Finetuning the Entire Network

Let's see what happens when we fine-tune the entire network. That is, instead of just learning the weights in the final layers, we'll let the weights of the entire network change during the training process.

```
In [15]: 1 model_p3 = torchvision.models.detection.maskrcnn_resnet50_fpn(pretrained=True)
```

#### Question 9 ¶

Just like in question 5, we're interested in fine-tuning the pretrained model for our new dataset, so we will again need to modify its box predictor (FastRCNNPredictor) and its mask predictor (MaskRCNNPredictor) to match with our new dataset. Complete the code cell below. Hint: This is exactly the same as question 5.

```
In [16]: 1 from torchvision.models.detection.faster_rcnn import FastRCNNPredictor
2 from torchvision.models.detection.mask_rcnn import MaskRCNNPredictor
3
4 # Our new dataset has 3 classes: butterfly, squirrel and background
5 num_classes = 3
6
7 # Get number of input features for the classifier
8 in_features = model_p3.roi_heads.box_predictor.cls_score.in_features
9
10 # TODO: replace the pre-trained head with a new one
11 model_p3.roi_heads.box_predictor = FastRCNNPredictor(in_features, num_classes)
12
13 # Get number of input features for the mask predictor
14 in_features_mask = model_p3.roi_heads.mask_predictor.conv5_mask.in_channels
15 hidden_layer = 256
16
17 # TODO: replace the mask predictor with a new one
18 model_p3.roi_heads.mask_predictor = MaskRCNNPredictor(in_features_mask, hidden_layer, num_classes)
```

Again, let's ensure that we're using the GPU by printing the device info. This finetuning process takes even longer because we have more trainable parameters, hence there will be more computations.

```
In [17]: 1 device = torch.device('cuda') if torch.cuda.is_available() else torch.device('cpu')
2 print("Using", device)
3
4 model_p3 = model_p3.to(device)
```

Using cuda

We'll use the same optimizer and learning rate scheduler as before.

```
In [18]: 1 params = [p for p in model_p3.parameters() if p.requires_grad]
2 optimizer = torch.optim.SGD(params, lr=0.005, momentum=0.9, weight_decay=0.0005)
3 lr_scheduler = torch.optim.lr_scheduler.StepLR(optimizer,
4                                                step_size=3,
5                                                gamma=0.1)
```

Let's start the finetuning process.

```
In [19]: 1 num_epochs = 3
2
3 from engine import train_one_epoch, evaluate
4
5 for epoch in range(num_epochs):
6     print("Epoch", epoch)
7     # train for one epoch, printing every 10 iterations
8     train_one_epoch(model_p3, optimizer, data_loader, device, epoch, print_freq=10)
9     # update the learning rate
10    lr_scheduler.step()
11    # evaluate on the test dataset
12    evaluate(model_p3, data_loader_test, device=device)
```

data: 0.0102 max mem: 4593  
Test: Total time: 0:00:27 (0.1361 s / it)  
Averaged stats: model\_time: 0.1249 (0.1225) evaluator\_time: 0.0034 (0.0030)  
Accumulating evaluation results...  
DONE (t=0.04s).  
Accumulating evaluation results...  
DONE (t=0.04s).  
IoU metric: bbox

Average Precision	(AP) @[ IoU=0.50:0.95   area= all   maxDets=100 ]	= 0.641
Average Precision	(AP) @[ IoU=0.50   area= all   maxDets=100 ]	= 0.960
Average Precision	(AP) @[ IoU=0.75   area= all   maxDets=100 ]	= 0.761
Average Precision	(AP) @[ IoU=0.50:0.95   area= small   maxDets=100 ]	= -1.000
Average Precision	(AP) @[ IoU=0.50:0.95   area=medium   maxDets=100 ]	= 0.589
Average Precision	(AP) @[ IoU=0.50:0.95   area= large   maxDets=100 ]	= 0.658
Average Recall	(AR) @[ IoU=0.50:0.95   area= all   maxDets= 1 ]	= 0.648
Average Recall	(AR) @[ IoU=0.50:0.95   area= all   maxDets= 10 ]	= 0.722
Average Recall	(AR) @[ IoU=0.50:0.95   area= all   maxDets=100 ]	= 0.722
Average Recall	(AR) @[ IoU=0.50:0.95   area= small   maxDets=100 ]	= -1.000
Average Recall	(AR) @[ IoU=0.50:0.95   area=medium   maxDets=100 ]	= 0.658
Average Recall	(AR) @[ IoU=0.50:0.95   area= large   maxDets=100 ]	= 0.736

We'll generate the output for the same image, but this time with the new finetuned model.

```
In [20]: 1 img, target = dataset_test[1]
2 type(img)
3 # put the model in evaluation mode
4 model_p3.eval()
5 with torch.no_grad():
6     prediction = model_p3([img.to(device)])
7
8 Image.fromarray(img.mul(255).permute(1, 2, 0).byte().numpy())
```



Here's the predicted image.

```
In [21]: 1 Image.fromarray(prediction[0]['masks'][0, 0].mul(255).byte().cpu().numpy())
```



**Question 10**

Does this model perform better than the trained model in part2? Explain why.

**Answer**

The mAP for the model in part2 is 53.7% and 58.9% for bounding box and mask, respectively. Whereas, in the last fine-tuning, we get mAP for the model as 64.1% and 73.5% for bbox and mask, respectively.

Thus, it can be said that this model performs better than the previous model.

**Part 4: Extra Credit Questions**

**Question 11 [15 points]**

Can fully convolutional networks (FCN) be used for object detection? In Mask-RCNN we have 3 branches — mask, classification, and bounding box regression — out of which the last 2 have fully connected (FC) layers. Can this entire pipeline be replaced by a fully convolutional network? If possible, give 1 or 2 networks to support your claim.

**Answer**

Yes, there are several studies that show that FCN can be used for end-to-end object detection.

Paper 1: "Object Detection by R-FCN: This is a region-based fully convolutional object detector. To achieve this goal, the authors proposes a position-sensitive score maps to address a dilemma between translation-invariance in image classification and translation-variance in object detection. Their method can thus naturally adopt fully convolutional image classifier backbones, such as the latest Residual Networks (ResNets) for object detection. They showed competitive results on the PASCAL VOC datasets with the 101-layer ResNet.

Link to above paper: <https://arxiv.org/abs/2012.03544> (<https://arxiv.org/abs/2012.03544>)

Paper 2: The authors adapted the pretrained networks (AlexNet, the VGG net, and GoogLeNet) into fully convolutional networks and transfer their learned representations by fine-tuning to the segmentation task. They, then defined a skip architecture that combines semantic information from a deep, coarse layer with appearance information from a shallow, fine layer to produce accurate and detailed segmentations. Their fully convolutional networks achieve improved segmentation of PASCAL VOC (30% relative improvement to 67.2% mean IU on 2012), NYUDv2, SIFT Flow, and PASCAL-Context, while inference took one tenth of a second for a typical image.

Link to above paper: E. Shelhamer, J. Long and T. Darrell, "Fully Convolutional Networks for Semantic Segmentation," in IEEE Transactions on Pattern Analysis and Machine Intelligence, vol. 39, no. 4, pp. 640-651, 1 April 2017, doi: 10.1109/TPAMI.2016.2572683.

**Question 12 [5 points]**

What is the advantage of using CONV layers over FC/Dense layers?

**Answer**

A mask encodes an input object’s spatial layout. Thus, unlike class labels or box offsets that are collapsed into short output vectors by fully-connected (fc) layers, extracting the spatial structure of masks can be addressed naturally by the pixel-to-pixel correspondence provided by convolutions.

Specifically, we predict the mask from each RoI using an FCN (Fully Conv Layers). This allows each layer in the mask branch to maintain the object spatial layout explicitly without collapsing it into a vector representation that lacks spatial dimensions.

Another advantage is, unlike previous methods that resort to fc layers for mask prediction, the fully convolutional representation requires fewer parameters, and is more accurate as demonstrated by experiments.



## Vision meets Machine Learning HW3, Part 2

Continuing from HW2 Part 2, We will train a face segmentation model using the mask generation function

```
In [1]: 1 import torch
2 import torchvision
3 import torch
4 import torchvision
5 import torch.nn as nn
6 import cv2
7 from collections import defaultdict
8 %matplotlib inline
9 import numpy as np
10 from matplotlib import pyplot as plt
```

/opt/conda/envs/py37/lib/python3.7/site-packages/tqdm/auto.py:22: TqdmWarning: IPProgress not found. Please update jupyter and ipywidgets. See [https://ipywidgets.readthedocs.io/en/stable/user\\_install.html](https://ipywidgets.readthedocs.io/en/stable/user_install.html) ([https://ipywidgets.readthedocs.io/en/stable/user\\_install.html](https://ipywidgets.readthedocs.io/en/stable/user_install.html))  
from .autonotebook import tqdm as notebook\_tqdm  
Unable to revert mtime: /opt/conda/fonts

```
In [2]: 1 mom = 0.9           # momemtum in batch normalization
2 lr = 0.0001          # Learning rate in model training
3 i_up_lim = 10        # the number of epochs
```

### Datagen

We will be using the 3500 images in shared/datasets/smaller\_faces

Generate masks for the first 3500 images from the smaller\_faces.txt

Update the paths in smaller\_faces.txt to the correct path (should point to the smaller\_faces folder in shared/datasets)

```
In [3]: 1 import cv2
2 import numpy as np
3 import math
4
5
6 #todo reuse hw2 code for datagen(Ellipse), This could take around 30-45 mins. make you use imwrite to save c
7
8 %run Mask.py
```

### Custom dataset

Create a custom dataset for this problem as we did in part1 of this HW

In [4]:

```
1 import os
2 import torch
3 import numpy as np
4 import matplotlib.pyplot as plt
5 from torch.utils.data import Dataset, DataLoader
6 from torchvision import transforms, utils
7 import glob
8 from PIL import Image
9
10 class face_segment(Dataset):
11
12     def __init__(self, x, y, x_transform = None, y_transform = None):
13         #todo
14         imgs = glob.glob(x + '/*.jpg')
15         self.imgs = sorted(imgs)
16         #print("The sorted imagelist is ", self.imgs)
17
18         masks = glob.glob(y + '/*.jpg')
19         self.masks = sorted(masks)
20         #print("The sorted maskslis is ", self.masks)
21
22         self.x_transform = x_transform
23         self.y_transform = y_transform
24
25     def __len__(self):
26         #todo
27         return len(self.imgs)
28
29     def __getitem__(self, idx):
30         #todo
31         #No fancy transforms needed. only the ones defined in the cell below.
32         #make sure the your masks are 1 channel and it is populated only with 0s or 1s
33
34         img_path = self.imgs[idx]
35         mask_path = self.masks[idx]
36
37         img = Image.open(img_path).convert("RGB")
38         mask = np.array(Image.open(mask_path))*255
39
40         if self.x_transform is not None:
41             img = self.x_transform(img)
42         if self.y_transform is not None:
43             mask = self.y_transform(mask)
44
45         return img, mask
```

In [5]:

```
1 import torchvision.transforms as T
2 def get_transform(X):
3     transforms = []
4     transforms.append(T.ToTensor())
5     if X:
6         transforms.append(T.Resize((800,800), interpolation=T.InterpolationMode.BILINEAR) )
7     else:
8         transforms.append(T.Resize((200,200), interpolation=T.InterpolationMode.NEAREST) )
9
10     return T.Compose(transforms)
```

Dataset and DataLoader

In [6]:

```
1 # use our dataset and defined transformations, it would look something like vv
2 # dataset = face_segment(X, Y, get_transform(train=True), get_transform(train=False))
3
4 X = "/home/jovyan/shared/datasets/smaller_faces"
5 Y = "/home/jovyan/HW3/Masks"
6 dataset = face_segment(X, Y, get_transform(True), get_transform(False))
7
8 # define training data loaders, it would look something like vv
9 # data_loader = torch.utils.data.DataLoader(dataset, batch_size=4, shuffle=True)
10
11 data_loader = torch.utils.data.DataLoader(dataset, batch_size=4, shuffle=True)
12
13 # idx=3000
14 # img, mask = dataset[idx]
15 # plt.imshow(img[0])
16 # #print(mask[0])
```

In [7]:

```
1 # plt.imshow(mask[0])
2 # print(mask.shape)
3 # print(img.shape)
4 # #print(img[0])
```

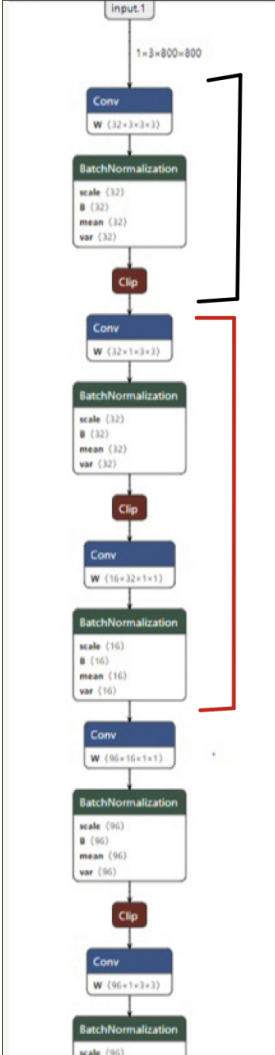
Centerface model

Hints: A lot of the code you need has already been put in comments below. please go through them and come back to this description for more clarity.

We will now be reconstructing a modified version of a model (modded\_centerface.onnx) described in <https://arxiv.org/abs/1911.03599> (<https://arxiv.org/abs/1911.03599>). Instead of building the entire detection pipeline, we will only create the classification branch (segmentation). This model uses mobilenet\_v2 as its backbone (encoder) and instead of coding the entire mobilenet\_v2 from scratch we will use mobilenet\_v2 from torch.hub.

Looking at the netron graph, you will realise that we need certain intermediate outputs of the mobilenet\_v2 model. How to get those outputs are described in the comments in the forward function.

You have to look at the at the outermost block numbers to get the value of which layer's output is added to which lane in the final block of the netron model as shown in this image. You can print the mobilenet\_v2 model and compare it with the netron graph as done below.



```
MobileNetV2(
  (features): Sequential(
    (0): ConvBNActivation(
      (0): Conv2d(3, 32, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
      (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (2): ReLU6(inplace=True)
    )
    (1): InvertedResidual(
      (conv): Sequential(
        (0): ConvBNActivation(
          (0): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), groups=32, bias=False)
          (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
          (2): ReLU6(inplace=True)
        )
        (1): Conv2d(32, 16, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (2): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      )
    )
    (2): InvertedResidual(
      (conv): Sequential(
        (0): ConvBNActivation(
          (0): Conv2d(16, 96, kernel_size=(1, 1), stride=(1, 1), bias=False)
          (1): BatchNorm2d(96, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
          (2): ReLU6(inplace=True)
        )
        (1): ConvBNActivation(
          (0): Conv2d(96, 96, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), groups=96, bias=False)
          (1): BatchNorm2d(96, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
          (2): ReLU6(inplace=True)
        )
        (2): Conv2d(96, 24, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (3): BatchNorm2d(24, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      )
    )
  )
)
```

In [8]:

```
1 import torch.nn as nn
2 from torch.nn import Sequential
3 import torch
4 import numpy as np
5 from torchvision import models
6 import torch.nn.functional as f
7 import cv2 as cv
8 from torchsummary import summary
9
10 class centerface(nn.Module):
11     def __init__(self):
12         super(centerface, self).__init__()
13         # These are all the layers you will need. You may use these as is or make the model yourself from s
14         # You have to fill out the forward function according to the netron graph
15
16         mobilenet_v2 = torch.hub.load('pytorch/vision:v0.10.0', 'mobilenet_v2', pretrained=True)
17         self.mobilenet_v2=nn.ModuleList(list(mobilenet_v2.features)[:1])
18
19         self.mobilenet_v2_final = Conv2d(320,24,1,1)
20         self.lane_0_transpose = ConvTranspose2d(24,24,2,2)
21         self.lane_1_conv = Conv2d(96,24,1,1)
22         self.lane_0_1_transpose = ConvTranspose2d(24,24,2,2)
23         self.lane_2_conv = Conv2d(32,24,1,1)
24         self.lane_0_1_2_transpose = ConvTranspose2d(24,24,2,2)
25         self.lane_3_conv = Conv2d(24,24,1,1)
26         self.final_merge = Conv2d(24,24,3,1,same_padding=True)
27         self.hmap = nn.Conv2d(24, 1, 1, 1, padding=0,bias=True)
28         self.sig_out=nn.Sigmoid()
29         self._initialize_weights()
30
31         # result = []
32         # for idx, model in enumerate(self.mobilenet_v2):
33         #     #print("idx is ", idx)
34         #     #print("model ", model)
35         #     if idx in [3,6,13,17]:
36         #         result.append(model)
37         #     print ("The result is as follows ")
38         #     print(result)
39
40     def _initialize_weights(self):
41         for m in self.modules():
42             if isinstance(m, nn.Conv2d):
43                 nn.init.normal_(m.weight, std=0.01)
44                 if m.bias is not None:
45                     nn.init.constant_(m.bias, 0)
46             elif isinstance(m, nn.BatchNorm2d):
47                 nn.init.constant_(m.weight, 1)
48                 nn.init.constant_(m.bias, 0)
49
50     def forward(self, x):
51
52         #todo
53         # You will need to get intermediate outputs from the mobilenet_v2 backbone. You can do that by itera
54         # over the nn.ModuleList(list(mobilenet_v2.features)[:1]) we created previously.
55         # ModuleList can be indexed like a regular Python list, but modules it contains are properly registe
56         # and will be visible by all Module methods.
57         # E.g,
58         # result=[]
59         # for idx,model in enumerate(self.mobilenet_v2):
60         #     x=model(x)
61         #     if(idx in {whichever layer indices you need (You will get this information from the netron gr
62         #         result.append(x)
63
64         result = []
65         for idx, model in enumerate(self.mobilenet_v2):
66             x =model(x)
67             if idx in [3,6,13,17]:
68                 result.append(x)
69
70         x_lane_0 = result[3]
71         x_lane_0 = self.mobilenet_v2_final(x_lane_0)
72         x_lane_0 = self.lane_0_transpose(x_lane_0)
73
74         x_lane_1 = result[2]
75         x_lane_1 = self.lane_1_conv(x_lane_1)
76
77         x1 = x_lane_0 + x_lane_1
78         x1 = self.lane_0_1_transpose(x1)
79
80         x_lane_2 = result[1]
81         x_lane_2 = self.lane_2_conv(x_lane_2)
82
83         x2 = x1 + x_lane_2
84         x2 = self.lane_0_1_2_transpose(x2)
85
86         x_lane_3 = result[0]
87         x3 = self.lane_3_conv(x_lane_3)
88
89         x_final = x2 + x3
90         x_final = self.final_merge(x_final)
```

```

91         x_final = self.hmap(x_final)
92         x = self.sig_out(x_final)
93
94         return x
95
96
97 class Conv2d(nn.Module):
98     def __init__(self, in_channels, out_channels, kernel_size, stride=1, relu=True, same_padding=False, bn=True):
99         super(Conv2d, self).__init__()
100         padding = int((kernel_size - 1) / 2) if same_padding else 0
101         self.conv = nn.Conv2d(in_channels, out_channels, kernel_size, stride, padding=padding,bias=False)
102         self.bn = nn.BatchNorm2d(out_channels, eps=0.001, momentum=mom, affine=True,track_running_stats=False)
103         self.relu = nn.ReLU(inplace=False) if relu else None
104
105     def forward(self, x):
106         x = self.conv(x)
107         if self.bn is not None:
108             x = self.bn(x)
109         if self.relu is not None:
110             x = self.relu(x)
111         return x
112
113 class ConvTranspose2d(nn.Module):
114     def __init__(self, in_channels, out_channels, kernel_size=2, stride=2, relu=True, same_padding=False, bn=True):
115         super(ConvTranspose2d, self).__init__()
116         padding = int((kernel_size - 1) / 2) if same_padding else 0
117         self.conv = nn.ConvTranspose2d(in_channels, out_channels, kernel_size, stride, padding=padding,bias=False)
118         self.bn = nn.BatchNorm2d(out_channels, eps=0.001, momentum=mom, affine=True,track_running_stats=False)
119         self.relu = nn.ReLU(inplace=False) if relu else None
120
121     def forward(self, x):
122         x = self.conv(x)
123         if self.bn is not None:
124             x = self.bn(x)
125         if self.relu is not None:
126             x = self.relu(x)
127         return x
128
129 #k = centerface()
130 #inputs = torch.randn(1,3,800,800,requires_grad=True)
131 #torch_out = k(inputs)
132 #summary(k,(3,200,200))
133 #torch.onnx.export(k,inputs,"My_CNN_model.onnx", export_params=True)
```

## Boilerplate utilities

In [9]:

```

1 import time
2 class AverageMeter(object):
3     """Computes and stores the average and current value"""
4     def __init__(self):
5         self.reset()
6
7     def reset(self):
8         self.val = 0
9         self.avg = 0
10        self.sum = 0
11        self.count = 0
12
13    def update(self, val, n=1):
14        self.val = val
15        self.sum += val * n
16        self.count += n
17        self.avg = self.sum / self.count
18
19 def log_print(text, color=None, on_color=None, attrs=None):
20     if cprint is not None:
21         cprint(text, color=color, on_color=on_color, attrs=attrs)
22     else:
23         print(text)
24 try:
25     from termcolor import cprint
26 except ImportError:
27     cprint = None
```

## Initialize model



In [10]:

```
1 #todo
2 model = centerface()
3 device = torch.device('cuda') if torch.cuda.is_available() else torch.device('cpu')
4 model = centerface().to(device).float()
```

Downloading: "https://github.com/pytorch/vision/archive/v0.10.0.zip" to /home/jovyan/.cache/torch/hub/v0.10.0.zip  
Downloading: "https://download.pytorch.org/models/mobilenet\_v2-b0353104.pth" to /home/jovyan/.cache/torch/hub/checkpoints/mobilenet\_v2-b0353104.pth  
100%|██████████| 13.6M/13.6M [00:00<00:00, 68.3MB/s]  
Using cache found in /home/jovyan/.cache/torch/hub/pytorch\_vision\_v0.10.0

# Training loop

You can take a look at this link to get an understanding of the training function.

<https://pytorch.org/tutorials/beginner/introyt/trainingyt.html> (<https://pytorch.org/tutorials/beginner/introyt/trainingyt.html>)

In [11]:

```
1 def train(data_loader, model, criterion, optimizer, epoch):
2     losses = AverageMeter()
3     batch_time = AverageMeter()
4     data_time = AverageMeter()
5     train_loader = data_loader
6
7     log_text = 'epoch %d, processed %d samples, lr %.10f' % (epoch, epoch * len(train_loader.dataset), lr)
8     # this is just for logs
9     log_print(log_text, color='green', attrs=['bold']) # this is just for logs
10
11     # todo
12     end = time.time()
13
14     # This is pseudocode.
15     for i, data in enumerate(train_loader):
16         imgs, masks = data
17         device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
18         imgs, masks = imgs.to(device), masks.to(device)
19
20         #I am filling up the boiler plate utilities just for consistent output logs.
21         data_time.update(time.time() - end) # this is just for logs
22         #todo...
23         #....
24
25         optimizer.zero_grad()
26         outputs = model(imgs)
27         loss = criterion(outputs, masks)
28         losses.update(float(loss), 1) # this is just for logs
29         loss.backward()
30         optimizer.step()
31
32         batch_time.update(time.time() - end)
33         end = time.time()
34
35         if i % 20 == 0:
36             log_text = (('Epoch: [{0}][{1}/{2}]\t'
37                         'Patch {patch_num:d}\t'
38                         'Time {batch_time.val:.3f} ({batch_time.avg:.3f})\t'
39                         'Data {data_time.val:.3f} ({data_time.avg:.3f})\t'
40                         'Loss {loss.val:.8f} ({loss.avg:.4f})\t'
41                         ).format(epoch, i, len(train_loader), patch_num=0, batch_time=batch_time,
42                                data_time=data_time, loss=losses))
43
44             log_print(log_text, color='green', attrs=['bold'])
45
46
```

# Train the model


```
In [12]: 1 optimizer = torch.optim.Adam(filter(lambda p: p.requires_grad, model.parameters()), lr,weight_decay=5e-3)
2 criterion = nn.BCELoss(reduction='mean').cuda()
3
4 for i in range(i_up_lim):
5     train(data_loader, model, criterion, optimizer, i)
6     torch.save(model, 'model.pt')
```

Epoch: [9][520/875]	Patch 0	Time 0.344 (0.339)	Data 0.149 (0.144)	Loss 0.06680539 (0.1007)
Epoch: [9][540/875]	Patch 0	Time 0.341 (0.339)	Data 0.145 (0.144)	Loss 0.08494115 (0.1004)
Epoch: [9][560/875]	Patch 0	Time 0.346 (0.339)	Data 0.153 (0.144)	Loss 0.08154360 (0.1007)
Epoch: [9][580/875]	Patch 0	Time 0.334 (0.339)	Data 0.132 (0.144)	Loss 0.10970800 (0.1005)
Epoch: [9][600/875]	Patch 0	Time 0.341 (0.339)	Data 0.146 (0.144)	Loss 0.30756873 (0.1013)
Epoch: [9][620/875]	Patch 0	Time 0.333 (0.339)	Data 0.138 (0.144)	Loss 0.09629402 (0.1011)
Epoch: [9][640/875]	Patch 0	Time 0.331 (0.339)	Data 0.131 (0.144)	Loss 0.06516825 (0.1012)
Epoch: [9][660/875]	Patch 0	Time 0.316 (0.339)	Data 0.122 (0.144)	Loss 0.05113391 (0.1013)
Epoch: [9][680/875]	Patch 0	Time 0.326 (0.339)	Data 0.131 (0.144)	Loss 0.15673082 (0.1015)
Epoch: [9][700/875]	Patch 0	Time 0.344 (0.339)	Data 0.149 (0.144)	Loss 0.07833005 (0.1017)
Epoch: [9][720/875]	Patch 0	Time 0.338 (0.339)	Data 0.142 (0.144)	Loss 0.06947359 (0.1017)
Epoch: [9][740/875]	Patch 0	Time 0.409 (0.339)	Data 0.215 (0.144)	Loss 0.06252707 (0.1015)
Epoch: [9][760/875]	Patch 0	Time 0.337 (0.339)	Data 0.141 (0.144)	Loss 0.10605373 (0.1019)
Epoch: [9][780/875]	Patch 0	Time 0.347 (0.339)	Data 0.154 (0.144)	Loss 0.13296019 (0.1026)
Epoch: [9][800/875]	Patch 0	Time 0.333 (0.339)	Data 0.140 (0.144)	Loss 0.06059437 (0.1028)
Epoch: [9][820/875]	Patch 0	Time 0.329 (0.339)	Data 0.131 (0.144)	Loss 0.04567843 (0.1033)
Epoch: [9][840/875]	Patch 0	Time 0.354 (0.339)	Data 0.160 (0.144)	Loss 0.17501679 (0.1034)
Epoch: [9][860/875]	Patch 0	Time 0.356 (0.339)	Data 0.161 (0.144)	Loss 0.08343604 (0.1030)

## Test the model and visualise the predictions

Pass the test image (test\_grid.jpg) to your model and visualise the model's prediction.

```
In [17]: 1 model.eval()
2
3 img = Image.open("/home/jovyan/HW3/test_grid.jpg")
4
5 transform = transforms.Compose([transforms.Resize((800,800)),transforms.ToTensor()])
6 input_img = transform(img)
7 input_img = input_img.unsqueeze(0)
8
9 if torch.cuda.is_available():
10     input_img = input_img.to('cuda')
11     model.to('cuda')
12
13 with torch.no_grad():
14     prediction = model(input_img)
15 prediction = prediction.detach().cpu()
16 prediction = np.squeeze(prediction)
17 plt.imshow(prediction)
18 plt.show()
```



```
In [ ]: 1
```

In [1]:

```
1 def Mask():
2     import cv2
3     import numpy as np
4     from IPython.display import Image
5     from PIL import Image as convert_to_image
6     import matplotlib.pyplot as plt
7     import os
8
9     replace_path = "/home/jovyan/shared/datasets/smaller_faces/"
10    source_path = "/home/jovyan/HW3/smaller_faces.txt"
11
12    with open(source_path, 'r') as file :
13        filedata = file.read()
14
15    filedata = filedata.replace("smaller_faces/", replace_path)
16
17    with open('smaller_faces_dup.txt', 'w') as file:
18        file.write(filedata)
19
20    ### -----#
21
22    def create_folder():
23        path = "/home/jovyan/HW3/Masks/"
24        isExist = os.path.exists(path)
25
26        if not isExist:
27            os.makedirs(path)
28        return path
29
30    ### -----#
31
32    def ellipse(box_coord, image_width, image_height, w_x, h_y):
33        image = np.zeros((image_height, image_width), dtype=int)
34        center = [box_coord[0]+w_x//2, box_coord[1]+h_y//2]
35        x, y = center[1], center[0] # Rows and columns of the center
36
37        cx,cy = [x, y]
38        queue = [[cx,cy]]
39        neighbors = [[1, 0], [-1, 0], [0, 1], [0, -1]]
40
41        if w_x < h_y:
42            min_ax = w_x//2
43            maj_ax = h_y//2
44        else:
45            min_ax = h_y//2
46            maj_ax = w_x//2
47
48        def dist(a,b):
49            return np.sqrt((((a[0]-b[0])**2)*(min_ax**2)) + (((a[1]-b[1])**2)*(maj_ax**2)))
50
51        while(len(queue)>0):
52            tx, ty = queue.pop(0)
53            if(dist([cx,cy], [tx,ty])) >= (min_ax*maj_ax) or image[ty,tx]==1:
54                continue # Do not need to fill this value, because outside the bounds
55            image[ty,tx] = 1 # mark the pixel; we'll multiply the mask by 255 for visualization later
56
57            for i,j in neighbors:
58                x = tx + i
59                y = ty + j
60                if(0<=x<img.shape[0] and 0<=y<img.shape[1]): #boundary check to make sure co-ordinates are
61                    queue.append((x,y))
62
63        return np.transpose(image)
64
65    ### -----#
66
67    def lookup():
68        heading_list = []
69        n_faces = []
70        dim_list = []
71        heading_name = []
72
73        with open("smaller_faces_dup.txt") as file:
74            lines = file.readlines()
75            file.close()
76
77        lookup = "/shared/datasets/smaller_faces/"
78
79        for num, line in enumerate(lines, 0):
80            if lookup in line:
81                heading_name.append(line.strip())
82                heading_list.append(int(num))
83                n = int(lines[num+1])
84                n_faces.append(n)
85                for i in range(n):
86                    dim_list.append(list(map(int, lines[num+2+i].split()[0:4])))
87
88        return heading_list, n_faces, dim_list, heading_name
89
90    ### -----#
```

